

No rewarded flow for AOL is ready to ship yet.

Nothing below is promoted from a weaker verdict: only SHIP proposals become flows.

We reviewed 3 ideas: 0 ship, 0 need revision, 3 rejected.

No idea cleared the judge's bar (weighted score ≥ 3.8 with every criterion ≥ 3 and every gate passed), so there is no flow to show. "Ideas we rejected, and why" lists what each one would need.

EXCHANGE RATE

Not computable


No priced packs observed: unit prices, action costs and exchange rate are unavailable.

CHEAPEST PACK

No priced pack was observed, so there is no paid path to cannibalize in currency terms.

How AOL makes money today

Monetizes primarily through native and display advertisements (such as Taboola advertising units) embedded within news feeds and article pages. Real screens from the app; prices and costs as observed.



PRICE PER UNIT
No priced packs observed

FREE EVERY DAY
No daily free source observed

CHEAPEST PACK
None observed

ADS TODAY
Banner and native sponsored recommendations (Taboola) appear frequently throughout the news feed and article view pages.

Ideas we rejected, and why

Every candidate the judge scored, with its final verdict. Only SHIP becomes a flow; REVISE is never promoted. Weighted score out of 5.

ID	IDEA	VERDICT	SCORE	WHY
P1	Unlock Commenting with Game Play Product change	REJECT	4.30	REJECT after 1 revision: revision stalled (weighted -0.2, below +0.2) with the same gate failures. Top concern: Grounding error in economy items ('POST_WALL' is not...
P3	Premium Insights Unlock Product change	REJECT	3.95	REJECT after 2 revisions: fixable gate failed: grounding (code); specificity scored 2 (< 3); still REVISE after round 2. Top concern: AOL is a traditional ad-supported...
P2	Daily News Reader Streak Product change	REJECT	3.65	REJECT after 1 revision: revision stalled (weighted -0.05, below +0.2) with the same gate failures. Top concern: Cannibalization of ad-free subscription value: giving...

Brainstormed, not developed: Ad-Free Reading Hour · Saved Articles Capacity · AI News Summary Boost · Offline Reading Refill · Sponsored Brand Topic

Obvious baseline the proposer had to go beyond: Insert more native recommendation units into the Home Feed feed. · Show full-screen interstitial ads between article clicks. · Add more persistent banner ads in article detail views.

Method and mock fidelity

- 1. **Explore.** 60 steps, 17 states, 58 edges, 1 external surface in 8 min; stopped: budget steps. 72 actions deliberately not explored.
- 2. **Understand.** 17 screens in the product model; 8 claims verified against on-screen text, 0 inferred. Synthesized by llm.
- 3. **Mock + QA.** 33% of screens rebuilt in HTML; 22/22 navigation flows pass.
- 4. **Propose.** 8 ideas, 3 developed into typed proposals (patch, storyboard, assumptions). Generated by llm.
- 5. **Judge.** 7 judgment rounds for 3 proposals: code gates first, then a rubric scored by llm; the verdict is computed in code (ship ≥ 3.8 , every criterion ≥ 3 , ≤ 2 rounds).
- 6. **Slides.** Every frame is a screenshot of the patched mock driven through its runtime API; callouts sit on measured DOM positions; numbers come from economics code.

QA FIDELITY (MOCK VS THE REAL SCREENSHOTS)

MEAN COMPOSITE		FLOW PASS RATE		HTML SCREENS	
0.77		100%		33%	
SCREEN	RENDER	COMPOSITE	ROUNDS	MUST-FIX	
Home News Feed	html	0.73	1	0	
Home News Feed	html	0.79	1	1	
Article Details	image	–	0	0	
Conversation Screen	image	–	0	0	
Add Comment Page	image	–	0	0	
Sign In Screen	image	–	0	0	
Home News Feed	html	0.65	1	1	
Home News Feed	image	–	0	0	
Saved Articles	image	–	0	0	
Search Suggestions	image	–	0	0	
Account Menu Sidebar	image	–	0	0	
No Contacts	image	–	0	0	

Cost and limitations

MODEL SPEND FOR THIS APP (FROM COST.JSONL)

STAGE	LIVE CALLS	CACHED	TOKENS IN	TOKENS OUT	USD
explore	29	0	114.0k	8.5k	\$0
understand	1	0	20.1k	1.4k	\$0
mock	7	0	56.7k	95.4k	\$0
qa	4	0	41.3k	61.1k	\$0
propose	8	0	404.7k	30.4k	\$0
judge	7	0	415.4k	18.9k	\$0
Total	56	0	1.1M	215.8k	\$0

Gemini free tier: calls cost \$0; tokens are still logged. Models: gemini-3.5-flash-lite, gemini-3.1-flash-lite. Cached calls replay from disk at no cost.

LIMITATIONS

- Explored on an Android emulator (logged-in account); 72 actions left unexplored by guard rails or scope, listed in the model viewer.
- 0 product-model claims inferred rather than observed; the economy review is the human checkpoint for them.
- Revenue per view uses KB eCPM ranges with a non-game haircut; some constants are inferred. ARPDAU appears only as a labelled scenario.
- The ad in the frames is simulated by the mock; real fill, creatives and completion rates vary by market.
- The judge is a separate call of the same model family as the proposer; its single-fault calibration table is in judge-eval.md.

How this runs across hundreds of apps

Storage and versioning

- Screenshots, element lists and HTML are content-addressed blobs; artifact JSON references them by hash.
- A model version is keyed by package, version code and a UI fingerprint of the core flows, which also catches over-the-air updates.
- A diff between two model versions (screens, walls, prices, offers) is itself a sales signal.

When to re-explore

- Daily store version poll; weekly no-LLM smoke replay of the core flows and the path to the wall.
- Drift above 10% or any price change triggers an incremental re-explore from the known graph (about 20% of first-run cost).
- A new ad format re-runs only propose and judge.

Where humans review

- App onboarding and login (15 min); economy review, the highest-leverage step (10 min).
- Deck approval before any customer sees it (15–30 min); exception queue for blocked apps.
- Monthly judge audit against labelled outcomes.

Cost and sales fit

- Prospect scan without a mock ≈ \$5 of model calls; full pitch pack ≈ \$20–25 plus 45–60 min of human review.
- The deck, clickable prototype and sizing sheet attach to the CRM deal; each SHIP's integration snippet becomes the implementation ticket.
- Post-launch opt-in, completion and IAP guardrails flow back to recalibrate economics and the judge.