


Let users **play a short game for a reward.**

Add one opt-in rewarded flow to AOL that trades a short sponsored game for something users already value. We reviewed 3 ideas: 1 ship, 0 need revision, 2 rejected.



1 FLOW 1 · SLIDE 3

Ad-Free for 5 Minutes for one short game

Enjoy 5 minutes of ad-free news reading after a quick game.

SHIP

Score 4.10 / 5 (ship ≥ 3.8)

Product change

\$0.009–\$0.015 per view

Earned per completed US view; the reward costs $\approx \$0.0063$ to serve.

Nothing free taken away

The reward is strictly time-boxed to 5 minutes and offered inline to non-payers, serving as a sampling effect for...

≤ 3 a day

Opt-in only, at least 1 h apart; nothing plays unless the user taps.

ONE COMPLETED VIEW EARNS

\$0.009–\$0.015


US rewarded eCPM after a non-game haircut, about \$0.0063 net at the low end. The app shows no prices, so there is no exchange rate in its own units; every reward on these slides costs less than that to serve. In LATAM a view earns \$0.0015–\$0.003.

CHEAPEST PACK

No priced pack was observed, so there is no paid path to cannibalize in currency terms.


How AOL makes money today

Monetizes primarily through native and display advertisements (such as Taboola advertising units) embedded within news feeds and article pages. Real screens from the app; prices and costs as observed.



Ad-Free for 5 Minutes for one short game

Enjoy 5 minutes of ad-free news reading after a quick game.



SHIP

Score 4.10 / 5 (ship ≥ 3.8)

Product change

WHY THIS WORKS

\$0.009–\$0.015 per view

Earned per completed US view; the reward costs $\approx$ \$0.0063 to serve.

Nothing free taken away

The reward is strictly time-boxed to 5 minutes and offered inline to non-payers, serving as a sampling effect for...

≤ 3 a day

Opt-in only, at least 1 h apart; nothing plays unless the user taps.

Open the clickable prototype →

Simula SIM-RWD · button entry · Game Partner:
AOL · min play 15s · reward granted on
REWARD_VERIFIED · frame callouts on slide 4

AOL Ad-Free Sprint

TRIGGER

User is browsing the News Feed Home or finishes reading an article.

ELIGIBILITY

Non-paying users currently exposed to standard native and banner ad density.

WHAT THE USER SEES

Ad-Free for 5 Minutes Play a 15-second game to remove ads for 5 minutes. Go Ad-Free with AOL Plus.

Go Ad-Free

Keep reading

REWARD AND CAPS

5 minutes of ad-free reading, 5 minutes; granted on REWARD_VERIFIED. At most 3 a day, 1 h apart.

FRAME BY FRAME (SLIDE 3)

- 1

Today The news feed displays integrated advertising units alongside the top stories.

User sees standard native ads in the feed.
- 2

What changed A value-exchange opportunity is introduced directly within the core reading loop.

New Ad-Free Sprint card appears inline.
- 3

Offer A Simula MiniGameInvitation explains the 15-second game and 5-minute reward.

Go Ad-Free: starts a 15s sponsored game

Keep reading: dismisses the offer

User taps to start the sprint.
- 4

Ad plays The user plays a quick 15-second game with the AOL character.
- 5

Value received Ads are removed from the feed for 5 minutes as the reward...

Timer begins and native ads are hidden.

ECONOMICS (COMPUTED IN CODE)

Revenue per completed US view	\$0.009–\$0.015
Reward value at list price	n/a (no priced unit)
Reward ÷ one view	n/a
Max earnable per day (3/day)	n/a
Cheapest pack	none observed
Cost to serve the reward	\$0.0063 per view

SCENARIO, NOT FORECAST

0.216

\$0.0026

\$946,080

views per DAU

ARPDau

per year per 1M DAU

Scenario, not a forecast: 12% of DAU engage, 1.8 views each, US mid eCPM after a 25% non-game haircut.

CANNIBALIZATION GUARD

The reward is strictly time-boxed to 5 minutes and offered inline to non-payers, serving as a sampling effect for potential ad-free subscription tiers.

KPI AND HOLDOUT

Net ARPU (total ad revenue per user); guardrails: Daily active news reading time, Native ad CTR in non-sprint sessions, D30 Retention. User-level holdout, 10% share, 28 days

RISKS

- The 5-minute duration must be monitored to ensure the rewarded eCPM exceeds the opportunity cost...
- User fatigue if the inline chip appears too frequently in short feeds.

INTEGRATION (@SIMULA/ADS-REACT-NATIVE)

```
import { SimulaAds, useRewardedAd, MiniGameButton } from "@simula/ads-react-native";
import { useEffect, useState } from "react";

// P2 on "News Feed Home". Trigger: User is browsing the News Feed Home (s13) or
// finishes reading an article.
// Caps 3/day, 60 min apart: app remote config (unit cap: dashboard)
const AD_UNIT = "SIM-RWD-news-feed-home"; // placeholder: create it in the dashboard

export function AoIAdFreeOffer({ userId, partnerImageUrl, grant, onDecline }) {
  const rwd = useRewardedAd(AD_UNIT);
  const [capped, setCapped] = useState(true);
  useEffect(() => {
    SimulaAds.checkFrequencyCap(AD_UNIT, userId).then(setCapped); // true: hide
    rwd.setMetadata({ surface: "news-feed-home", proposal: "P2" });
    rwd.load({ charName: "AOL" }); // Game Partner
  }, [userId]);
  useEffect(() => {
    // Grant on REWARD_VERIFIED (server-verified), never on EARNED_REWARD.
    if (rwd.rewardVerified) grant({ amount: 5, resource: "Ad-free Sprint", duration:
"5 minutes", token: rwd.rewardToken });
  }, [rwd.rewardVerified]);
  // no_fill or not loaded: render nothing; today's path is unchanged.
  if (capped || !rwd.isLoaded || rwd.error?.code === "no_fill") return null;
  // "Keep reading" stays the app's own control: back to where the user was.
  return (
    <MiniGameButton
      text={"Go Ad-Free"}
      onClick={() => rwd.show()}
    />
  );
}
// Min play 15s is a unit setting (no RN load option; unverified).
// Needs <SimulaProvider apiKey={...}> at the root; grant() is your server call.
```

WHAT THE JUDGE CHANGED

- Round 0 (v1) REVISE 3.30: The opportunity cost of suppressing 15 minutes of high-density display and native ads (\$0.0188) exceeds the revenue... Required: Shorten the...
- Round 1 (v2) REVISE 4.45: The unit economics are currently negative; 10 minutes of ad-free reading likely suppresses more ad revenue than a... Required: Reduce the...

PRECEDENTS (KNOWLEDGE BASE)

- TAX-2** Time-boxed premium unlock or trial ("sponsored session")
- EX-MUSIC** Spotify and Pandora (audio; freemium subscription)
- CORE-2**
- TRIG-1** Surface types, ranked by intent (from high to low)

Ideas we rejected, and why

Every candidate the judge scored, with its final verdict. Only SHIP becomes a flow; REVISE is never promoted. Weighted score out of 5.

ID	IDEA	VERDICT	SCORE	WHY
P2	AOL Ad-Free Sprint Product change	SHIP	4.10	SHIP at 4.1 (v3 after 2 revisions). See slide 3.
P3	Daily Reader Streak & Ad-Free Access Product change	REJECT	3.85	REJECT after 1 revision: revision stalled (weighted -0.15, below +0.2) with the same gate failures. Top concern: economics (code): 15 ad-free minutes give up about 15...
P1	AOL Guest Commenter Product change	REJECT	3.40	REJECT after 1 revision: revision stalled (weighted -1.2, below +0.2) with the same gate failures. Top concern: grounding (code): economy item "Account sign-in /...

Brainstormed, not developed: Headline Doubler · Saved List Expansion · Inbox Priority Pass · Sports Spotlight Session · Comment sample fallback

Obvious baseline the proposer had to go beyond: Watch a video to remove banner ads for 30 minutes of news reading. · Watch an ad to unlock a 'Premium' news article otherwise hidden behind a wall. · Watch a video to earn 100 generic coins for a new profile customization shop.

Method and mock fidelity

- 1. **Explore.** 60 steps, 17 states, 58 edges, 1 external surface in 8 min; stopped: budget steps. 72 actions deliberately not explored.
- 2. **Understand.** 17 screens in the product model; 8 claims verified against on-screen text, 0 inferred. Synthesized by llm.
- 3. **Mock + QA.** 33% of screens rebuilt in HTML; 22/22 navigation flows pass.
- 4. **Propose.** 8 ideas, 3 developed into typed proposals (patch, storyboard, assumptions). Generated by llm.
- 5. **Judge.** 7 judgment rounds for 3 proposals: code gates first, then a rubric scored by llm; the verdict is computed in code (ship ≥ 3.8 , every criterion ≥ 3 , ≤ 2 rounds).
- 6. **Slides.** Every frame is a screenshot of the patched mock driven through its runtime API; callouts sit on measured DOM positions; numbers come from economics code.

QA FIDELITY (MOCK VS THE REAL SCREENSHOTS)

MEAN COMPOSITE		FLOW PASS RATE		HTML SCREENS	
0.77		100%		33%	
SCREEN	RENDER	COMPOSITE	ROUNDS	MUST-FIX	
Home News Feed	html	0.73	1	0	
Home News Feed	html	0.79	1	1	
Article Details	image	–	0	0	
Conversation Screen	image	–	0	0	
Add Comment Page	image	–	0	0	
Sign In Screen	image	–	0	0	
Home News Feed	html	0.65	1	1	
Home News Feed	image	–	0	0	
Saved Articles	image	–	0	0	
Search Suggestions	image	–	0	0	
Account Menu Sidebar	image	–	0	0	
No Contacts	image	–	0	0	

Cost and limitations

MODEL SPEND FOR THIS APP (FROM COST.JSONL)

STAGE	LIVE CALLS	CACHED	TOKENS IN	TOKENS OUT	USD
explore	29	0	114.0k	8.5k	\$0
understand	1	0	20.1k	1.4k	\$0
mock	7	0	56.7k	95.4k	\$0
qa	4	0	41.3k	61.1k	\$0
propose	28	17	1.5M	108.1k	\$0
judge	30	9	1.6M	91.3k	\$0
slides	4	8	34.6k	33.8k	\$0
Total	103	34	3.3M	399.6k	\$0

Gemini free tier: calls cost \$0; tokens are still logged. Models: gemini-3.5-flash-lite, gemini-3.1-flash-lite, gemini-3-flash-preview, gemini-2.5-flash, gemini-3.7-flash, gemini-3.5-flash. Cached calls replay from disk at no cost.

LIMITATIONS

- Explored on an Android emulator (logged-in account); 72 actions left unexplored by guard rails or scope, listed in the model viewer.
- 0 product-model claims inferred rather than observed; the economy review is the human checkpoint for them.
- Revenue per view uses KB eCPM ranges with a non-game haircut; some constants are inferred. ARPDau appears only as a labelled scenario.
- The ad in the frames is simulated by the mock; real fill, creatives and completion rates vary by market.
- The judge is a separate call of the same model family as the proposer; its single-fault calibration table is in judge-eval.md.

How this runs across hundreds of apps

Storage and versioning

- Screenshots, element lists and HTML are content-addressed blobs; artifact JSON references them by hash.
- A model version is keyed by package, version code and a UI fingerprint of the core flows, which also catches over-the-air updates.
- A diff between two model versions (screens, walls, prices, offers) is itself a sales signal.

When to re-explore

- Daily store version poll; weekly no-LLM smoke replay of the core flows and the path to the wall.
- Drift above 10% or any price change triggers an incremental re-explore from the known graph (about 20% of first-run cost).
- A new ad format re-runs only propose and judge.

Where humans review

- App onboarding and login (15 min); economy review, the highest-leverage step (10 min).
- Deck approval before any customer sees it (15–30 min); exception queue for blocked apps.
- Monthly judge audit against labelled outcomes.

Cost and sales fit

- Prospect scan without a mock ≈ \$5 of model calls; full pitch pack ≈ \$20–25 plus 45–60 min of human review.
- The deck, clickable prototype and sizing sheet attach to the CRM deal; each SHIP's integration snippet becomes the implementation ticket.
- Post-launch opt-in, completion and IAP guardrails flow back to recalibrate economics and the judge.