

No rewarded flow for Janitor is ready to ship yet.

Nothing below is promoted from a weaker verdict: only SHIP proposals become flows.

We reviewed 3 ideas: 0 ship, 0 need revision, 3 rejected.

No idea cleared the judge's bar (weighted score ≥ 3.8 with every criterion ≥ 3 and every gate passed), so there is no flow to show. "Ideas we rejected, and why" lists what each one would need.

ONE COMPLETED VIEW EARNS

\$0.009–\$0.015

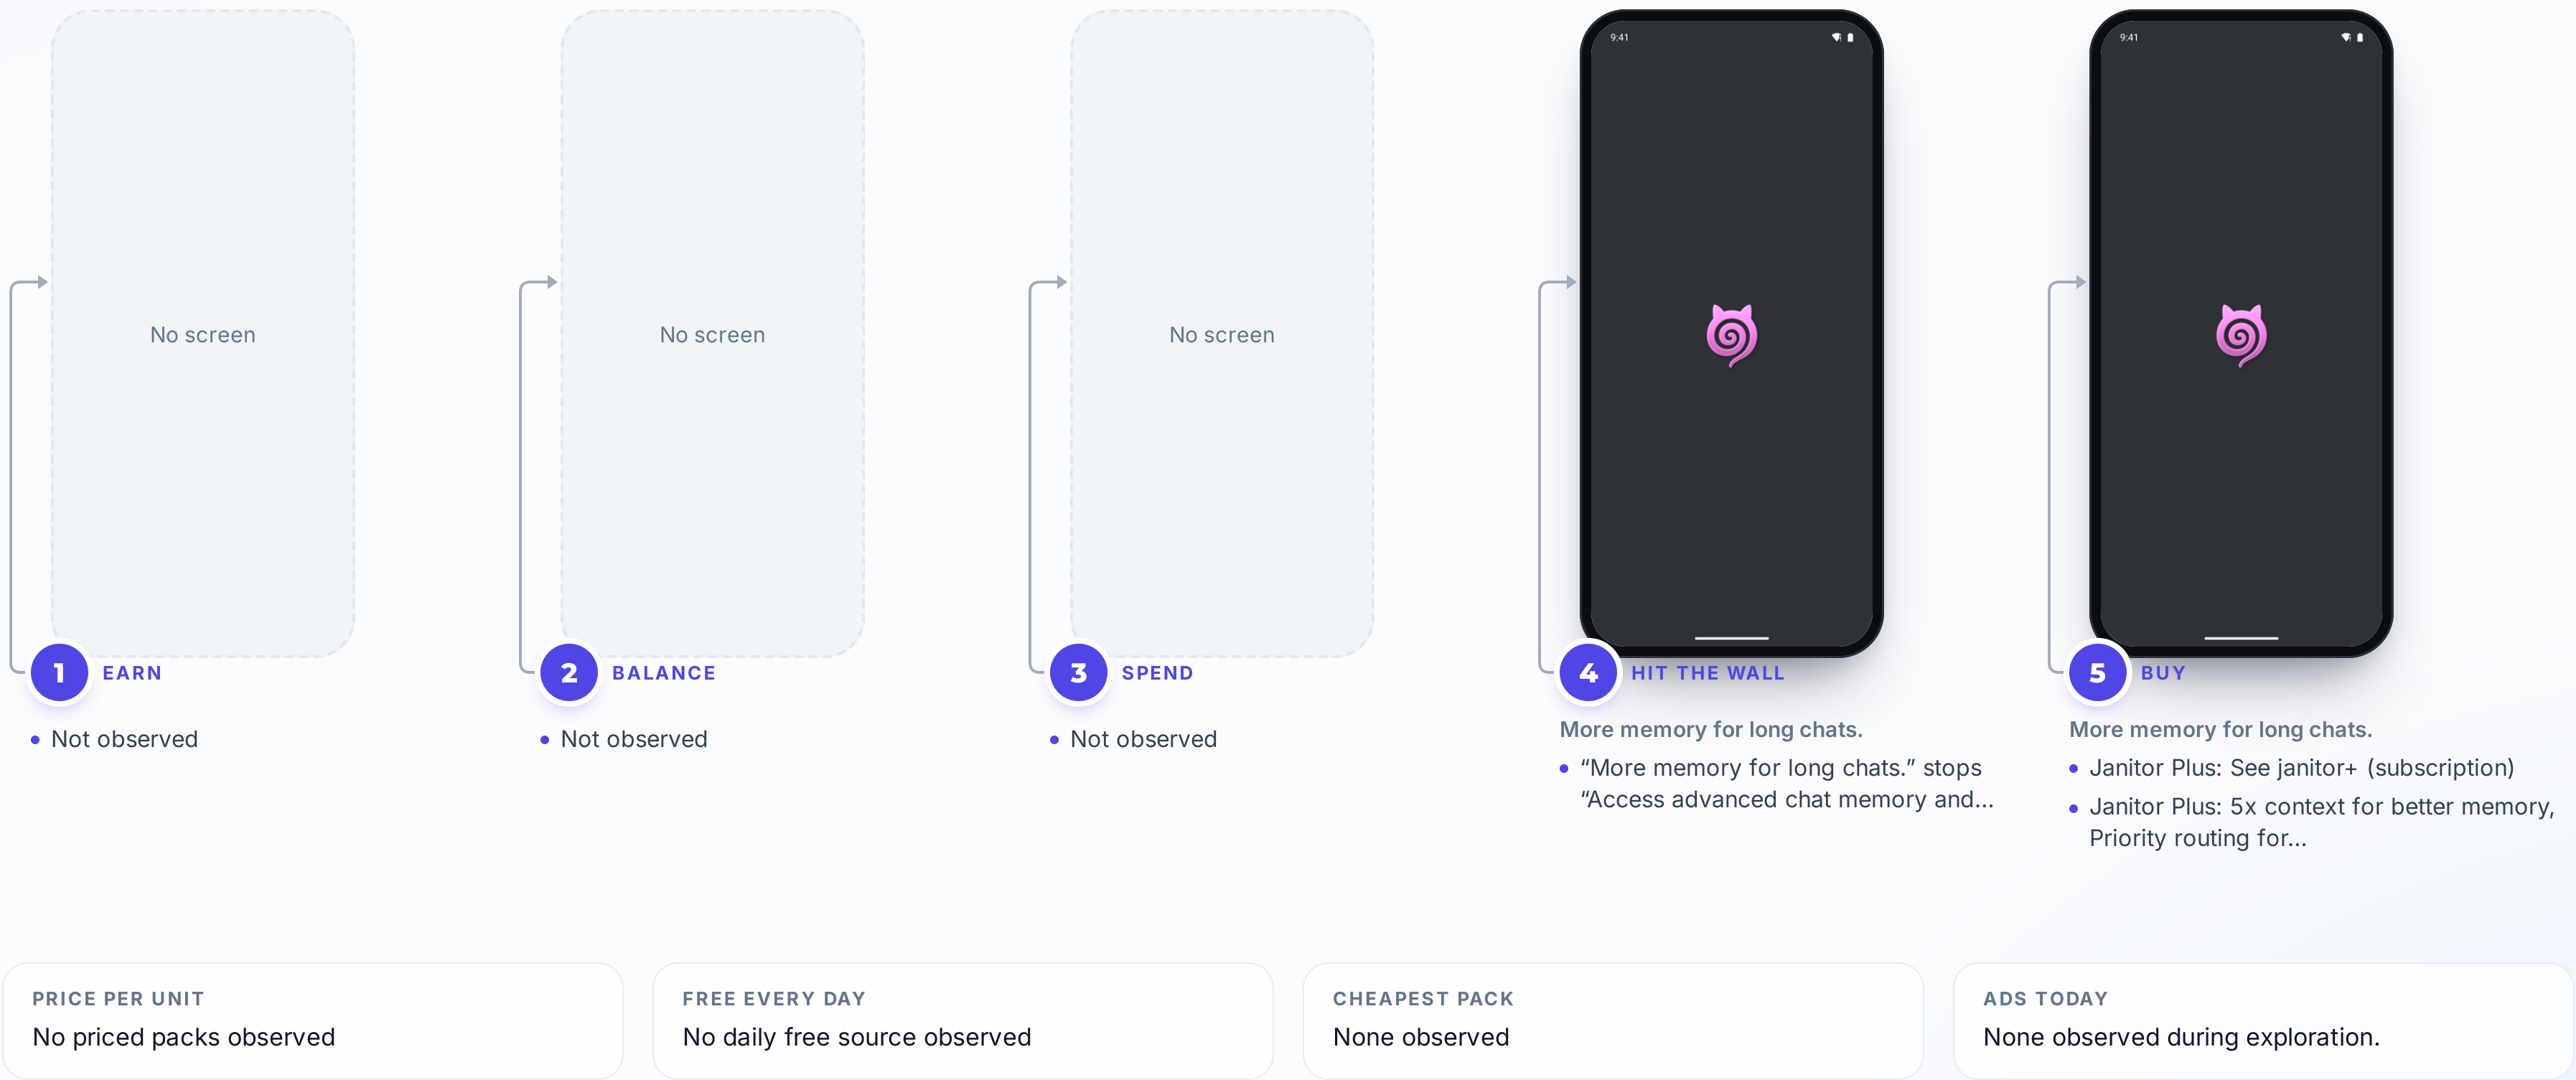
US rewarded eCPM after a non-game haircut, about \$0.0063 net at the low end. The app shows no prices, so there is no exchange rate in its own units; every reward on these slides costs less than that to serve. In LATAM a view earns \$0.0015–\$0.003.

CHEAPEST PACK

No priced pack was observed, so there is no paid path to cannibalize in currency terms.

How Janitor makes money today

Through the Janitor Plus subscription tier, which unlocks enhanced memory context, priority routing, monthly model swipes, and a golden username badge. Real screens from the app; prices and costs as observed.



Ideas we rejected, and why

Every candidate the judge scored, with its final verdict. Only SHIP becomes a flow; REVISE is never promoted. Weighted score out of 5.

ID	IDEA	VERDICT	SCORE	WHY
P1	Paywall Fallback Swipes Existing mechanic	REJECT	4.80	REJECT after 1 revision: policy gate failed: sfw (llm). Top concern: grounding (code): economy item "plan" does not exist; economy item "base_swipes (new)" does not...
P3	Daily Check-in Swipe Product change	REJECT	4.35	REJECT after 1 revision: revision stalled (weighted +0.05, below +0.2) with the same gate failures. Top concern: economics (code): Cost to serve the reward (\$0.0090)...
P2	Session Context Boost Existing mechanic	REJECT	4.15	REJECT after 1 revision: revision stalled (weighted -0.35, below +0.2) with the same gate failures. Top concern: economics (code): Cost to serve the reward (\$0.0090)...

Brainstormed, not developed: Timed Priority Routing Pass · Character Pre-Chat Priority Pass · Extended Character Persona Slot · Scene Memory Pin Refill · Profile Menu Frontier Swipes Refill

Obvious baseline the proposer had to go beyond: Watch a 30-second video ad to unlock 24 hours of free Janitor Plus subscription access. · Watch an ad when reaching a message limit to receive 10 more free messages. · Watch a rewarded ad from the profile menu to earn generic roleplay tokens.

Method and mock fidelity

- 1. **Explore.** 60 steps, 16 states, 60 edges, 0 external surfaces in 5 min; stopped: budget steps. 82 actions deliberately not explored.
- 2. **Understand.** 16 screens in the product model; 3 claims backed by a quote found on screen (evidence, not proof), 0 without one. Synthesized by llm.
- 3. **Mock + QA.** 33% of screens rebuilt in HTML; 45/45 navigation flows pass.
- 4. **Propose.** 8 ideas, 3 developed into typed proposals (patch, storyboard, assumptions). Generated by llm.
- 5. **Judge.** 6 judgment rounds for 3 proposals: code gates first, then a rubric scored by llm; the verdict is computed in code (ship ≥ 3.8 , every criterion ≥ 3 , ≤ 2 rounds).
- 6. **Slides.** Every frame is a screenshot of the patched mock driven through its runtime API; callouts sit on measured DOM positions; numbers come from economics code.

QA FIDELITY (MOCK VS THE REAL SCREENSHOTS)

MEAN COMPOSITE

0.90

FLOW PASS RATE

100%

HTML SCREENS

33%

SCREEN	RENDER	COMPOSITE	ROUNDS	MUST-FIX
Build, Share, Explore	html	0.94	1	0
@your_handle	html	0.92	0	0
More memory for long chats.	html	0.91	1	1
Janitor Plus Paywall	html	0.83	1	1
Explore Characters	image	–	0	0
Search Page	image	–	0	0
Search Characters	image	–	0	0
Search Characters	image	–	0	0
Character Details	image	–	0	0
Character Detail Page	image	–	0	0
Search Characters	image	–	0	0
Search Characters	image	–	0	0

Cost and limitations

MODEL SPEND FOR THIS APP (FROM COST.JSONL)

STAGE	LIVE CALLS	CACHED	TOKENS IN	TOKENS OUT	USD
explore	15	0	49.1k	3.1k	\$0
understand	1	0	20.2k	1.2k	\$0
mock	6	0	35.7k	69.8k	\$0
qa	3	0	13.5k	22.9k	\$0
propose	28	14	1.5M	121.1k	\$0
judge	38	25	2.1M	115.4k	\$0
slides	6	14	48.1k	61.6k	\$0
Total	97	53	3.8M	395.2k	\$0

Gemini free tier: calls cost \$0; tokens are still logged. Models: gemini-3.5-flash-lite, gemini-3.1-flash-lite, gemini-3-flash-preview, gemini-3.5-flash, gemini-3.8-flash, gemini-3.6-flash, gemini-2.5-flash. Cached calls replay from disk at no cost.

LIMITATIONS

- Explored on an Android emulator (logged-in account); 82 actions left unexplored by guard rails or scope, listed in the model viewer.
- 0 product-model claims without an on-screen quote, and a found quote is evidence, not proof; the economy review is the human checkpoint for both.
- Revenue per view uses KB eCPM ranges with a non-game haircut; some constants are inferred. ARPDau appears only as a labelled scenario.
- The ad in the frames is simulated by the mock; real fill, creatives and completion rates vary by market.
- The judge is a separate call of the same model family as the proposer; its single-fault calibration table is in judge-eval.md.

How this runs across hundreds of apps

Storage and versioning

- Screenshots, element lists and HTML are content-addressed blobs; artifact JSON references them by hash.
- A model version is keyed by package, version code and a UI fingerprint of the core flows, which also catches over-the-air updates.
- A diff between two model versions (screens, walls, prices, offers) is itself a sales signal.

When to re-explore

- Daily store version poll; weekly no-LLM smoke replay of the core flows and the path to the wall.
- Drift above 10% or any price change triggers an incremental re-explore from the known graph (about 20% of first-run cost).
- A new ad format re-runs only propose and judge.

Where humans review

- App onboarding and login (15 min); economy review, the highest-leverage step (10 min).
- Deck approval before any customer sees it (15–30 min); exception queue for blocked apps.
- Monthly judge audit against labelled outcomes.

Cost and sales fit

- Prospect scan without a mock ≈ \$5 of model calls; full pitch pack ≈ \$20–25 plus 45–60 min of human review.
- The deck, clickable prototype and sizing sheet attach to the CRM deal; each SHIP's integration snippet becomes the implementation ticket.
- Post-launch opt-in, completion and IAP guardrails flow back to recalibrate economics and the judge.