



Let users play a short game for actions.

Add two opt-in rewarded flows to Luzia that trade a short sponsored game for actions users already spend. We reviewed 5 ideas: 2 ship, 0 need revision, 3 rejected.



1 FLOW 1 · SLIDE 3

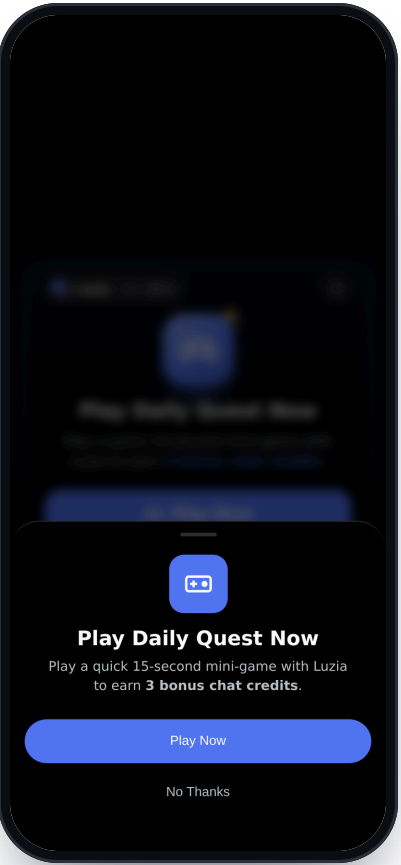
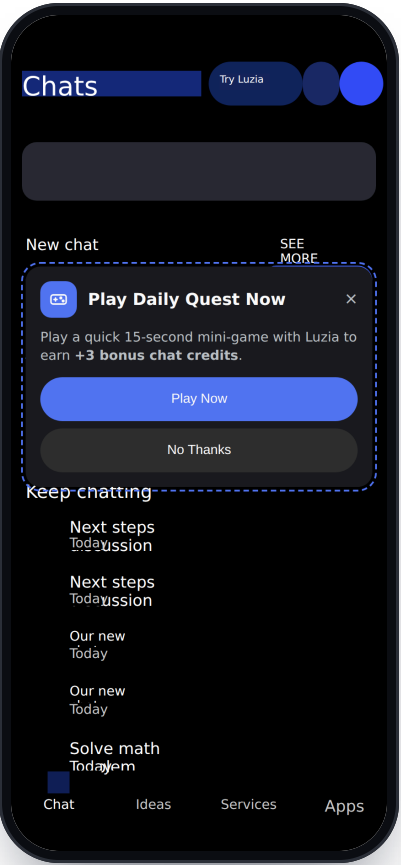
Create Account Sheet: play for 1 task access

Unlock immediate access to premium task services by playing a quick mini-game.

SHIP

Score 4.75 / 5 (ship ≥ 3.8)

Existing mechanic



2 FLOW 2 · SLIDE 5

Chats Home: play for 3 bonus chat credits

Earn bonus chat credits daily by completing a quick interactive mini-game with Luzia.

SHIP

Score 4.45 / 5 (ship ≥ 3.8)

Product change

EXCHANGE RATE

Not computable

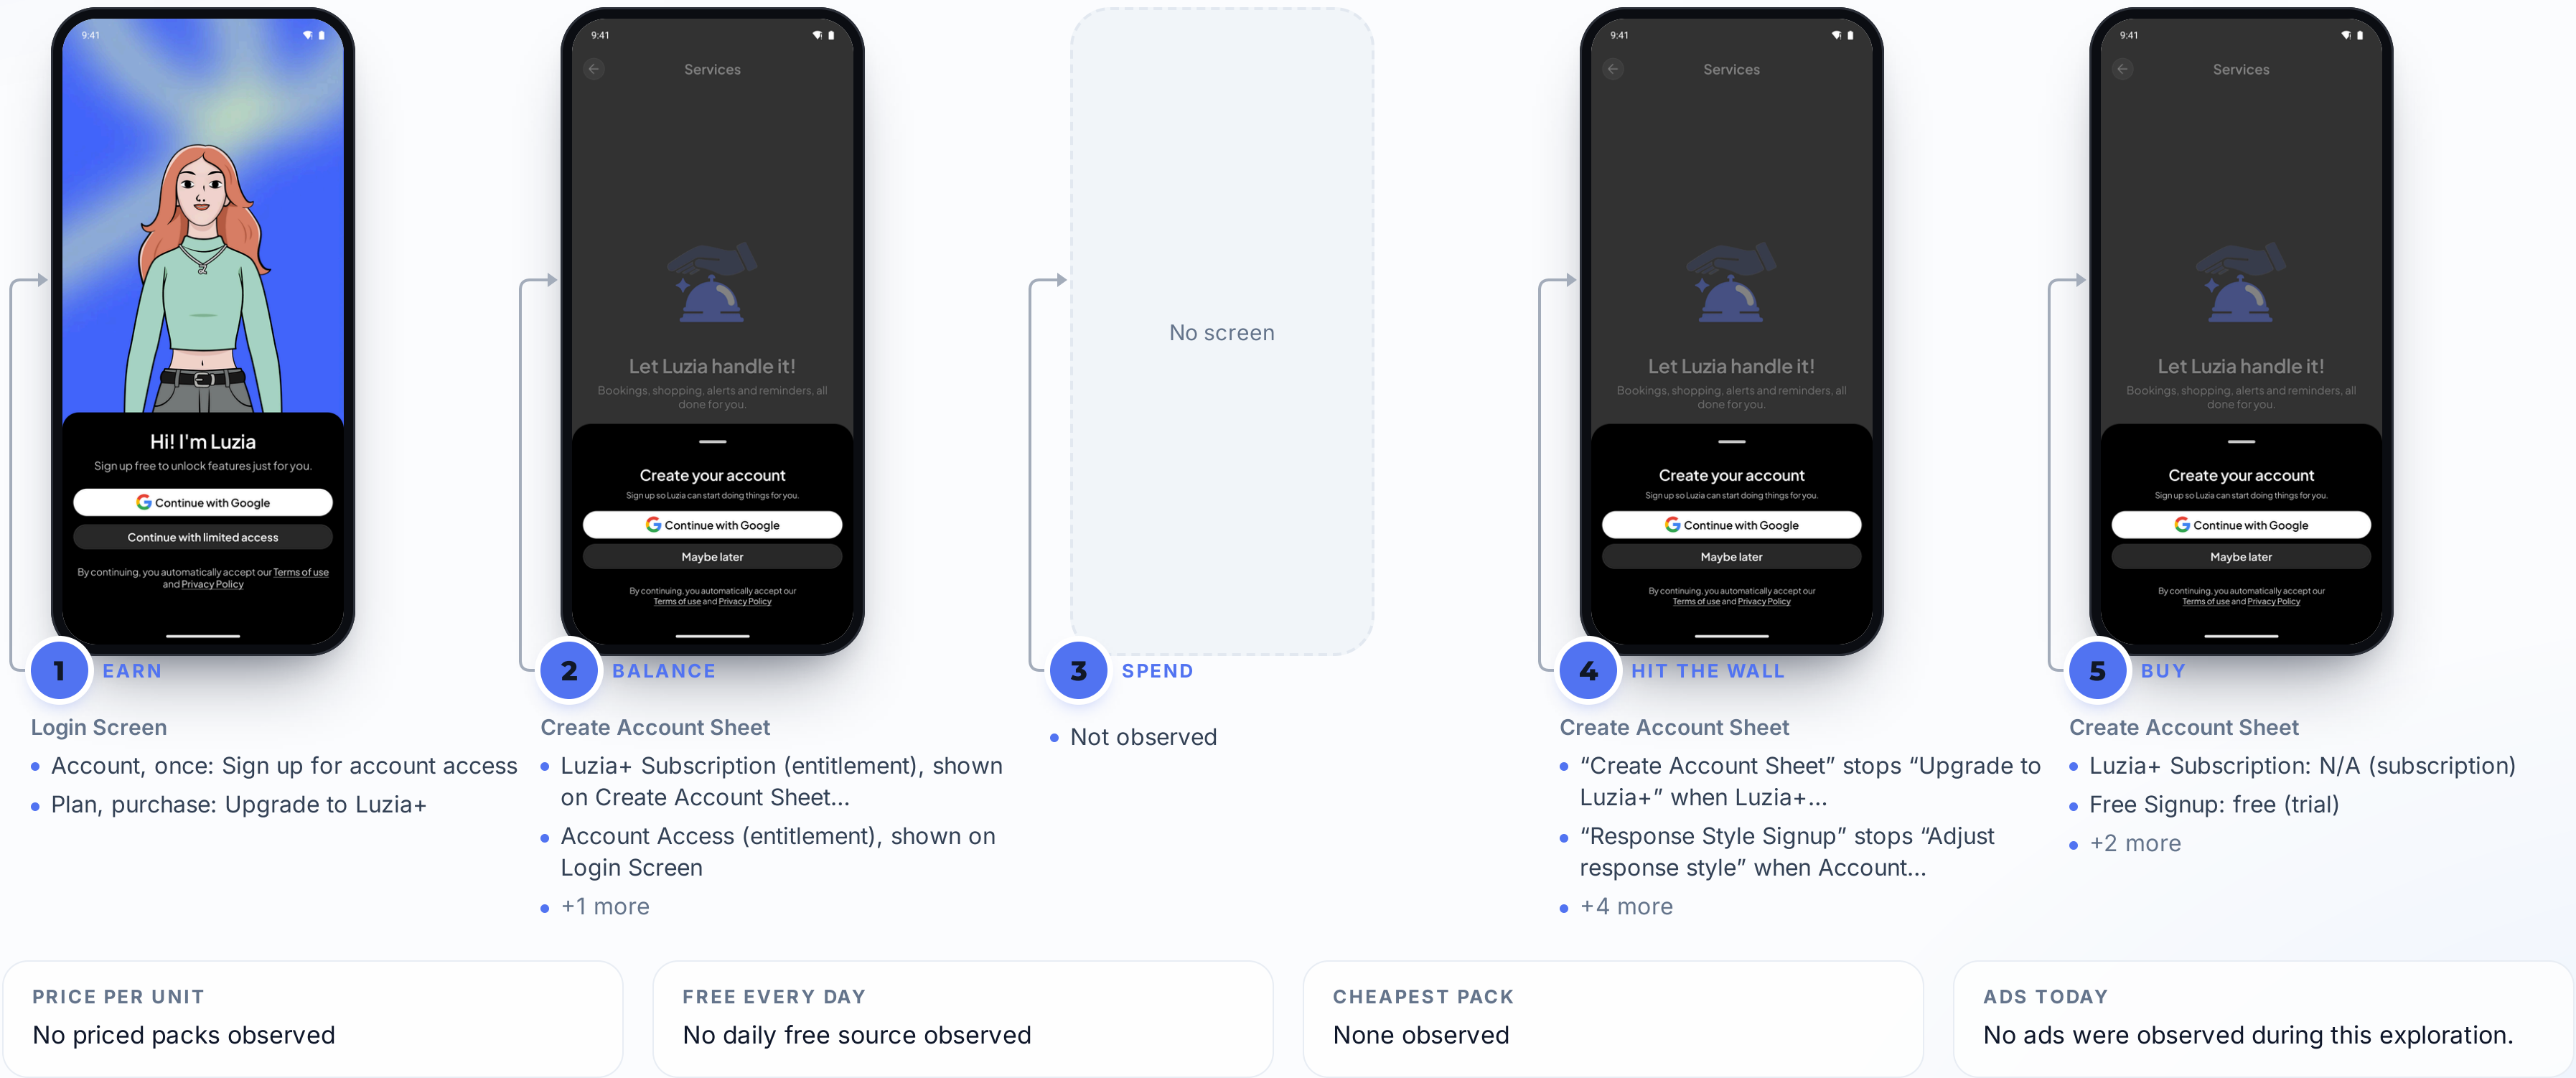
No priced packs observed: unit prices, action costs and exchange rate are unavailable.

CHEAPEST PACK

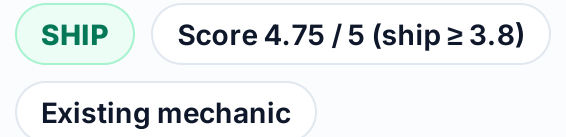
No priced pack was observed, so there is no paid path to cannibalize in currency terms.

How Luzia makes money today

The app monetizes through a premium subscription model ('Luzia+') that unlocks advanced features, and through a free signup that provides access to core customization... Real screens from the app; prices and costs as observed.



TRIGGER User attempts to start a task on the Services Tab and is blocked by the subscription wall



\$0.009–\$0.015 per view

Paid path untouched

Limited to 2 completions per day,
keeping usage far below the value of the
full Luzia+ subscription.

≤ 2 a day

Opt-in only, at least 30 min apart;
nothing plays unless the user taps.

Open the clickable prototype →

Simula SIM-RWD · button entry · Game Partner:
Luzia · min play 15s · reward granted on
REWARD_VERIFIED · frame callouts on slide 4

Start Task with Rewarded Refill

TRIGGER

User attempts to start a task on the Services Tab and is blocked by the subscription wall.

ELIGIBILITY

Non-paying users hitting the subscription wall on s07.

WHAT THE USER SEES

Start Task Now Play a quick 15-second game with Luzia to unlock this task.

Play Now

No thanks

REWARD AND CAPS

1 task access; granted on REWARD_VERIFIED. At most 2 a day, 30 min apart.

FRAME BY FRAME (SLIDE 3)

- 1

Today User wants to start a task.

A

User taps start new task
- 2

What changed Subscription wall blocks the action.

NEW

Subscription wall appears
- 3

Offer User is offered rewarded unlock.

B

Play Now: starts a 15s sponsored game

C

No thanks: dismisses the offer

Play now to unlock
- 4

Ad plays User plays mini-game.

Mini-game with Luzia
- 5

Value received Task is successfully unlocked.

D

Task unlocked

ECONOMICS (COMPUTED IN CODE)

Revenue per completed US view	\$0.009–\$0.015
Reward value at list price	n/a (no priced unit)
Reward ÷ one view	n/a
Max earnable per day (2/day)	n/a
Cheapest pack	none observed
Cost to serve the reward	\$0.0009 per view

SCENARIO, NOT FORECAST

0.6

\$0.0072

\$2,628,000

views per DAU

ARPDau

per year per 1M DAU

Scenario, not a forecast: 40% of DAU engage, 1.5 views each, US mid eCPM after a 25% non-game haircut.

CANNIBALIZATION GUARD

Limited to 2 completions per day, keeping usage far below the value of the full Luzia+ subscription. The offer appears only at the wall, ensuring the subscription remains the superior, unlimited path.

KPI AND HOLDOUT

Daily tasks completed by non-subscribers; guardrails: Luzia+ subscription conversion rate, Retention (D7, D30), Ad-related churn rate. 10% of non-subscribing users will not see the reward offer for 4 weeks.

RISKS

- Users may prefer ad-watching over subscribing if the task rewards are too frequent.
- Brand safety concerns if ads are not strictly SFw.

INTEGRATION (@SIMULA/ADS-REACT-NATIVE)

```
import { SimulaAds, useRewardedAd, MiniGameButton } from "@simula/ads-react-native";
import { useEffect, useState } from "react";

// P1 on "Create Account Sheet". Trigger: User attempts to start a task on the
// Services Tab and is blocked by the subscription wall.
// Caps 2/day, 30 min apart: app remote config (unit cap: dashboard)
const AD_UNIT = "SIM-RWD-create-account-sheet"; // placeholder: create it in the
// dashboard

export function StartTaskRewardedOffer({ userId, partnerImageUrl, grant, onDecline })
{
  const rwd = useRewardedAd(AD_UNIT);
  const [capped, setCapped] = useState(true);
  useEffect(() => {
    SimulaAds.checkFrequencyCap(AD_UNIT, userId).then(setCapped); // true: hide
    rwd.setMetadata({ surface: "create-account-sheet", proposal: "P1" });
    rwd.load({ charName: "Luzia" }); // Game Partner
  }, [userId]);
  useEffect(() => {
    // Grant on REWARD_VERIFIED (server-verified), never on EARNED_REWARD.
    if (rwd.rewardVerified) grant({ amount: 1, resource: "actions", token:
rwd.rewardToken });
  }, [rwd.rewardVerified]);
  // no_fill or not loaded: render nothing; today's path is unchanged.
  if (capped || !rwd.isLoaded || rwd.error?.code === "no_fill") return null;
  // "No thanks" stays the app's own control: back to where the user was.
  return (
    <MiniGameButton
      text={"Play Now"}
      onClick={() => rwd.show()}
    />
  );
}
// Min play 15s is a unit setting (no RN load option; unverified).
// Needs <SimulaProvider apiKey={...}> at the root; grant() is your server call.
```

WHAT THE JUDGE CHANGED

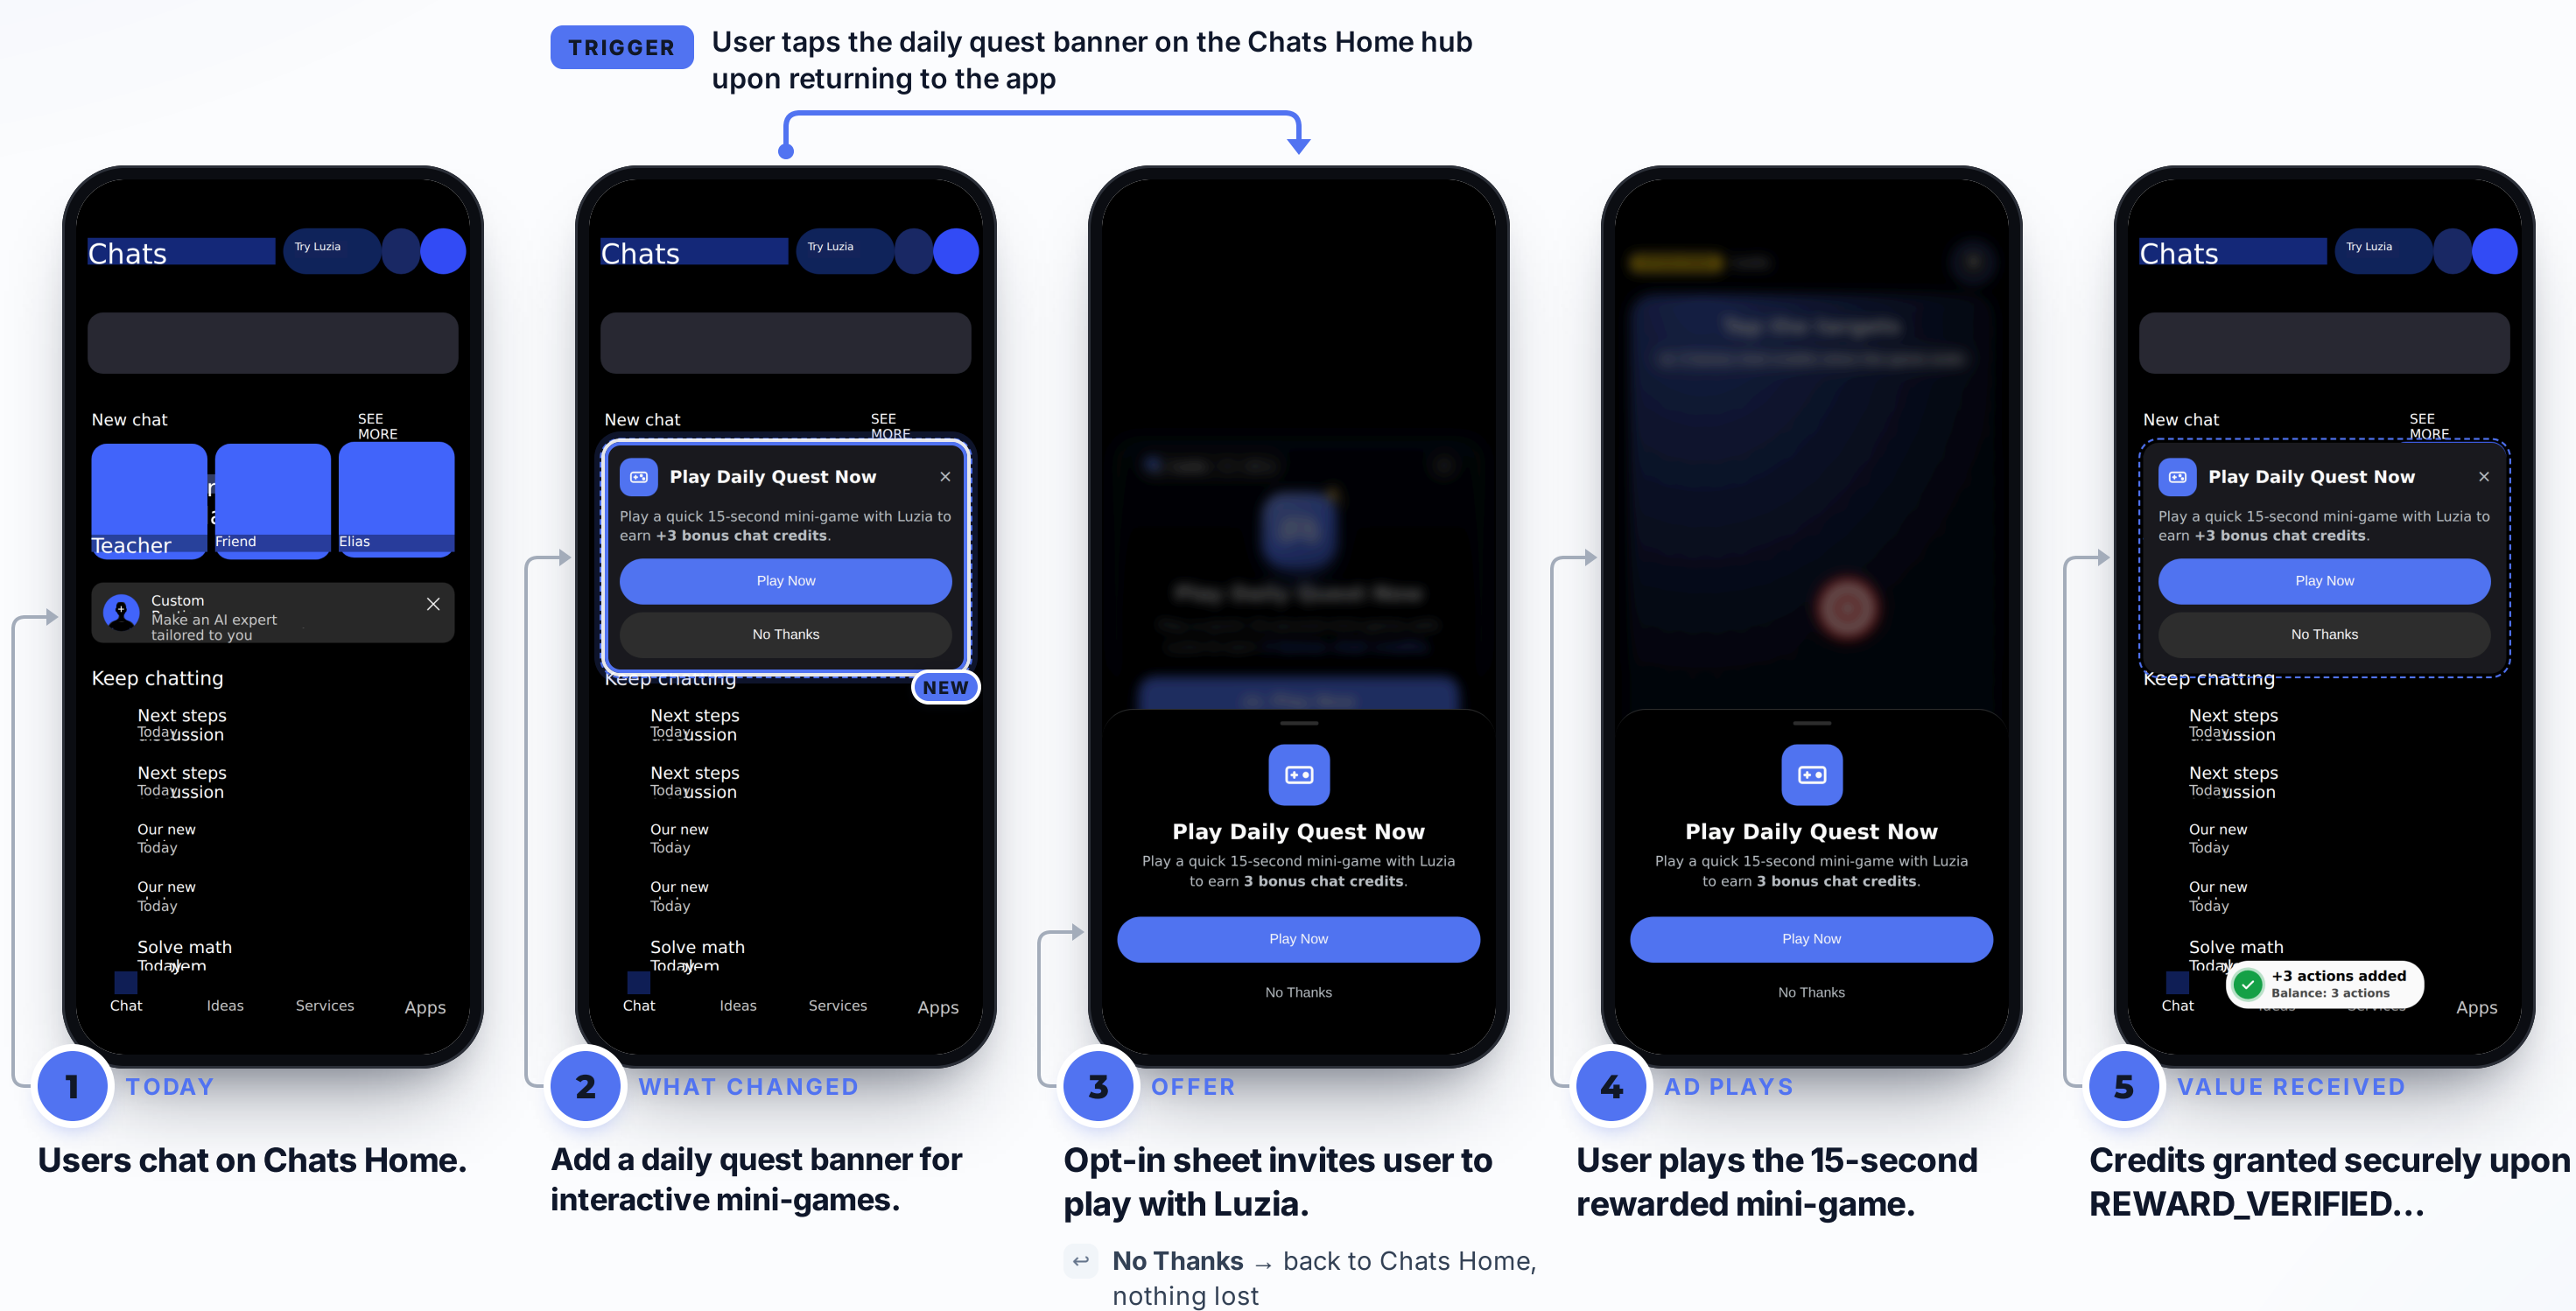
- Round 0 (v1) REVISE 3.80: Unit economics show that the cost to serve the task (\$0.0090) exceeds net revenue per view (\$0.0063) at the low end... Required: Fix element...

PRECEDENTS (KNOWLEDGE BASE)

- TAX-1** Consumable refill (messages, credits, energy, lives, puzzles)
- EX-DUO** Duolingo (education; subscription plus ads plus gems)
- AI-1**

Chats Home: play for 3 bonus chat credits

Earn bonus chat credits daily by completing a quick interactive mini-game with Luzia.



SHIP

Score 4.45 / 5 (ship ≥ 3.8)

Product change

WHY THIS WORKS

\$0.009–\$0.015 per view

Earned per completed US view; the reward costs ≈ \$0.0054 to serve.

Paid path untouched

Gated to non-paying users with strict daily caps, rewarding soft usage quota rather than permanent subscription...

≤ 3 a day

Opt-in only, at least 1 h apart; nothing plays unless the user taps.

Open the clickable prototype →

Simula SIM-RWD · button entry · Game Partner: Luzia · min play 15s · reward granted on REWARD_VERIFIED · frame callouts on slide 6

Daily Character Quest & Credits

TRIGGER

User taps the daily quest banner on the Chats Home hub upon returning to the app.

ELIGIBILITY

Non-paying users who have completed their first session and have remaining daily quota capacity.

WHAT THE USER SEES

Play Daily Quest Now Play a quick 15-second mini-game with Luzia to earn 3 bonus chat credits.

Play Now

No Thanks

REWARD AND CAPS

3 bonus chat credits; granted on REWARD_VERIFIED. At most 3 a day, 1 h apart.

FRAME BY FRAME (SLIDE 5)

- 1

Today Users chat on Chats Home without daily task rewards today.

Luzia chat home hub
- 2

What changed Add a daily quest banner for interactive mini-games.

NEW

Daily Quest banner
- 3

Offer Opt-in sheet invites user to play with Luzia.

Play Now: starts a 15s sponsored gameNo Thanks: dismisses the offerPlay 15 seconds
- 4

Ad plays User plays the 15-second rewarded mini-game.

Mini-game active
- 5

Value received Credits granted securely upon REWARD_VERIFIED callback.

+3 credits credited

ECONOMICS (COMPUTED IN CODE)

Revenue per completed US view	\$0.009–\$0.015
Reward value at list price	n/a (no priced unit)
Reward ÷ one view	n/a
Max earnable per day (3/day)	n/a
Cheapest pack	none observed
Cost to serve the reward	\$0.0054 per view

SCENARIO, NOT FORECAST

0.375

\$0.0045

\$1,642,500

views per DAU

ARPDau

per year per 1M DAU

Scenario, not a forecast: 25% of DAU engage, 1.5 views each, US mid eCPM after a 25% non-game haircut.

CANNIBALIZATION GUARD

Gated to non-paying users with strict daily caps, rewarding soft usage quota rather than permanent subscription features or premium reasoning models.

KPI AND HOLDOUT

Rewarded ARPDau and D7/D30 user retention.; guardrails: Paid conversion non-inferiority (~3% max relative drop), Daily message volume per active user, App store.... User-level 10% randomized holdout running for 28 days.

RISKS

- Low risk of minor fatigue if daily caps are exceeded
- Ensure game partner tone remains fully SFW and aligned with Luzia's brand

INTEGRATION (@SIMULA/ADS-REACT-NATIVE)

```
import { SimulaAds, useRewardedAd, MiniGameButton } from "@simula/ads-react-native";
import { useEffect, useState } from "react";

// P4 on "Chats Home". Trigger: User taps the daily quest banner on the Chats Home
// hub upon returning to the app.
// Caps 3/day, 60 min apart: app remote config (unit cap: dashboard)
const AD_UNIT = "SIM-RWD-chats-home"; // placeholder: create it in the dashboard

export function DailyCharacterQuestOffer({ userId, partnerImageUrl, grant, onDecline }) {
  const rwd = useRewardedAd(AD_UNIT);
  const [capped, setCapped] = useState(true);
  useEffect(() => {
    SimulaAds.checkFrequencyCap(AD_UNIT, userId).then(setCapped); // true: hide
    rwd.setMetadata({ surface: "chats-home", proposal: "P4" });
    rwd.load({ charName: "Luzia" }); // Game Partner
  }, [userId]);
  useEffect(() => {
    // Grant on REWARD_VERIFIED (server-verified), never on EARNED_REWARD.
    if (rwd.rewardVerified) grant({ amount: 3, resource: "actions", token:
rwd.rewardToken });
  }, [rwd.rewardVerified]);
  // no_fill or not loaded: render nothing; today's path is unchanged.
  if (capped || !rwd.isLoaded || rwd.error?.code === "no_fill") return null;
  // "No Thanks" stays the app's own control: back to where the user was.
  return (
    <MiniGameButton
      text={"Play Now"}
      onClick={() => rwd.show()}
    />
  );
}

// Min play 15s is a unit setting (no RN load option; unverified).
// Needs <SimulaProvider apiKey={...}> at the root; grant() is your server call.
```

WHAT THE JUDGE CHANGED

- Round 0 (v1) REVISE 4.50: The cost to serve the reward (\$0.0180) exceeds net revenue per view (\$0.0063) at the low end, resulting in negative... Required: Reduce the...

PRECEDENTS (KNOWLEDGE BASE)

- TAX-9** Daily quests, tasks and check-ins (sponsored tasks)
- AI-3**
- EX-DUO** Duolingo (education; subscription plus ads plus gems)

Ideas we rejected, and why

Every candidate the judge scored, with its final verdict. Only SHIP becomes a flow; REVISE is never promoted. Weighted score out of 5.

ID	IDEA	VERDICT	SCORE	WHY
P1	Start Task with Rewarded Refill Existing mechanic	SHIP	4.75	SHIP at 4.75 (v2 after 1 revision). See slide 3.
P4	Daily Character Quest & Credits Product change	SHIP	4.45	SHIP at 4.45 (v2 after 1 revision). See slide 5.
P5	Attachment Studio Extra Credit Product change	REJECT	4.50	REJECT after 1 revision: revision stalled (weighted +0.1, below +0.2) with the same gate failures. Top concern: Image generation COGS (\$0.0250) exceeds net revenue per...
P3	Deep Reasoning Session Pass Product change	REJECT	3.85	REJECT after 2 revisions: fixable gate failed: economics (code); unit-economics scored 1 (< 3); still REVISE after round 2. Top concern: The cost of the reward...
P2	AI Usage Quota Refill Fallback Existing mechanic	REJECT	3.65	REJECT after 2 revisions: fixable gate failed: economics (code); unit-economics scored 2 (< 3); weighted 3.65 < 3.8; still REVISE after round 2. Top concern: The cost...

Brainstormed, not developed: Response Style Trial Pass · Custom Bestie Sample Creation · Services Hub Daily Check-in Bonus · Task Queue Priority Boost · Settings Hub Ad-Free Session

Obvious baseline the proposer had to go beyond: Watch an ad to unlock Luzia+ subscription · Watch a video ad to get free messages · Watch an ad to remove all ads

Method and mock fidelity

- 1. **Explore.** 175 steps, 38 states, 171 edges, 1 external surface in 26 min; stopped: budget steps. 107 actions deliberately not explored.
- 2. **Understand.** 38 screens in the product model; 15 claims verified against on-screen text, 0 inferred. Synthesized by llm.
- 3. **Mock + QA.** 33% of screens rebuilt in HTML; 134/135 navigation flows pass.
- 4. **Propose.** 10 ideas, 5 developed into typed proposals (patch, storyboard, assumptions). Generated by llm.
- 5. **Judge.** 12 judgment rounds for 5 proposals: code gates first, then a rubric scored by llm; the verdict is computed in code (ship ≥ 3.8 , every criterion ≥ 3 , ≤ 2 rounds).
- 6. **Slides.** Every frame is a screenshot of the patched mock driven through its runtime API; callouts sit on measured DOM positions; numbers come from economics code.

QA FIDELITY (MOCK VS THE REAL SCREENSHOTS)

MEAN COMPOSITE		FLOW PASS RATE		HTML SCREENS	
0.72		99%		33%	
SCREEN	RENDER	COMPOSITE	ROUNDS	MUST-FIX	
Chats Home	html	0.80	1	5	
Custom Bestie Signup Sheet	image	–	0	0	
Ideas Feed	html	0.62	1	10	
Ideas Feed	html	0.79	1	0	
Photo Idea Sheet	html	0.57	1	0	
Services Tab	html	0.86	1	4	
Create Account Sheet	image	–	0	0	
Apps Catalog	html	0.83	1	0	
Animate Tool Screen	html	0.77	1	0	
Image Creation Hub	html	0.57	1	0	
Photoedit Screen	image	–	0	0	
Attach Photos	image	–	0	0	

Cost and limitations

MODEL SPEND FOR THIS APP (FROM COST.JSONL)

STAGE	LIVE CALLS	CACHED	TOKENS IN	TOKENS OUT	USD
explore	108	23	306.2k	62.8k	\$0
understand	3	0	58.1k	21.9k	\$0
mock	11	0	70.6k	97.7k	\$0
qa	11	0	71.8k	135.1k	\$0
judge	34	0	2.5M	103.8k	\$0
propose	18	0	1.1M	71.7k	\$0
slides	6	6	41.3k	35.5k	\$0
Total	191	29	4.1M	528.4k	\$0

Gemini free tier: calls cost \$0; tokens are still logged. Models: gemini-3.5-flash-lite, gemini-3.8-flash, gemini-2.5-flash, gemini-3.5-flash, gemini-3-flash-preview, gemini-3.1-flash-lite. Cached calls replay from disk at no cost.

LIMITATIONS

- Explored on an Android emulator (guest account); 107 actions left unexplored by guard rails or scope, listed in the model viewer.
- 0 product-model claims inferred rather than observed; the economy review is the human checkpoint for them.
- Revenue per view uses KB eCPM ranges with a non-game haircut; some constants are inferred. ARPDau appears only as a labelled scenario.
- The ad in the frames is simulated by the mock; real fill, creatives and completion rates vary by market.
- The judge is a separate call of the same model family as the proposer; its single-fault calibration table is in judge-eval.md.

How this runs across hundreds of apps

Storage and versioning

- Screenshots, element lists and HTML are content-addressed blobs; artifact JSON references them by hash.
- A model version is keyed by package, version code and a UI fingerprint of the core flows, which also catches over-the-air updates.
- A diff between two model versions (screens, walls, prices, offers) is itself a sales signal.

When to re-explore

- Daily store version poll; weekly no-LLM smoke replay of the core flows and the path to the wall.
- Drift above 10% or any price change triggers an incremental re-explore from the known graph (about 20% of first-run cost).
- A new ad format re-runs only propose and judge.

Where humans review

- App onboarding and login (15 min); economy review, the highest-leverage step (10 min).
- Deck approval before any customer sees it (15–30 min); exception queue for blocked apps.
- Monthly judge audit against labelled outcomes.

Cost and sales fit

- Prospect scan without a mock ≈ \$5 of model calls; full pitch pack ≈ \$20–25 plus 45–60 min of human review.
- The deck, clickable prototype and sizing sheet attach to the CRM deal; each SHIP's integration snippet becomes the implementation ticket.
- Post-launch opt-in, completion and IAP guardrails flow back to recalibrate economics and the judge.